

TinyLlama

TinyLlama

1. Model scale
 2. Pull TinyLlama
 3. Use TinyLlama
 - 3.1. Run TinyLlama
 - 3.2. Have a conversation
 - 3.3. End the conversation
- References

Demo Environment

Development board: Jetson Orin series motherboard

SSD: 128G

Tutorial application scope: Whether the motherboard can run is related to the available memory of the system. The user's own environment and the programs running in the background may cause the model to fail to run.

| Motherboard model | Run directly with Ollama | Run with Open WebUI |
|----------------------|--------------------------|---------------------|
| Jetson Orin NX 16GB | √ | √ |
| Jetson Orin NX 8GB | √ | √ |
| Jetson Orin Nano 8GB | √ | √ |
| Jetson Orin Nano 4GB | √ | √ |


1. Model scale

| Model | Parameters |
|-----------|------------|
| TinyLlama | 1.1B |

2. Pull TinyLlama

Using the pull command will automatically pull the model from the Ollama model library:

```
ollama pull tinyllama:1.1b
```



3. Use TinyLlama

3.1. Run TinyLlama

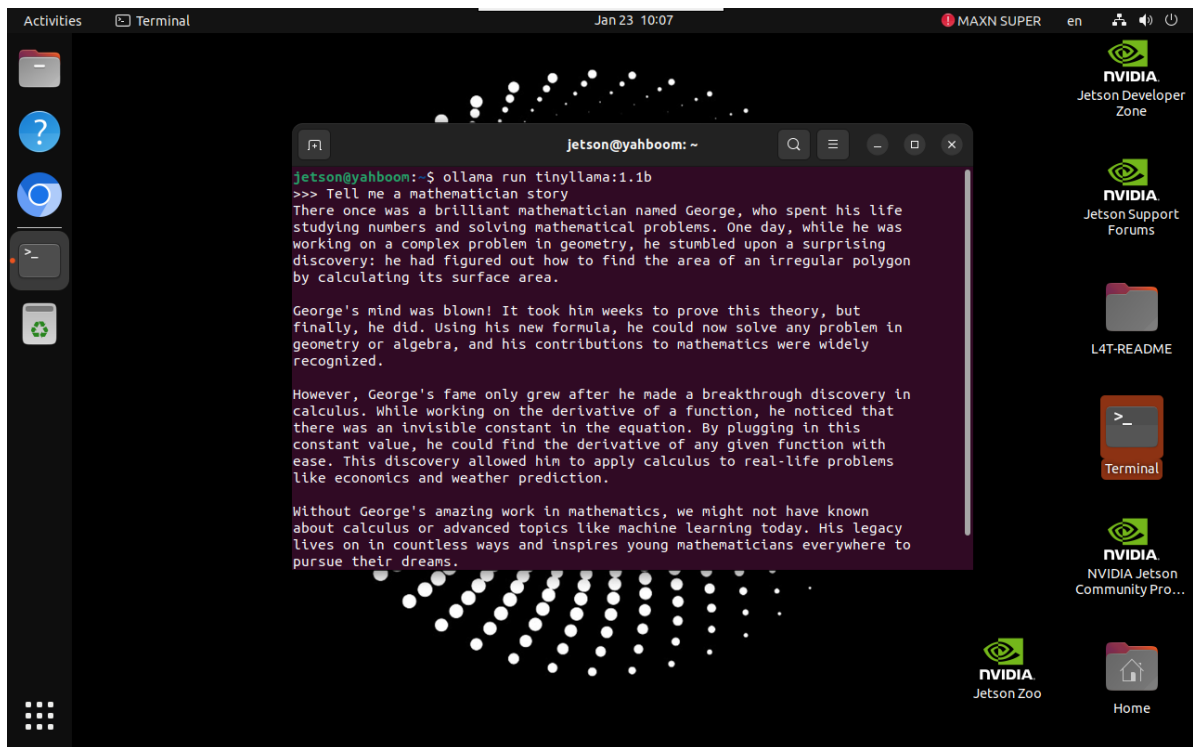
If the system does not have a running model, the system will automatically pull the TinyLlama 1.1B model and run it:

```
ollama run tinyllama:1.1b
```

3.2. Have a conversation

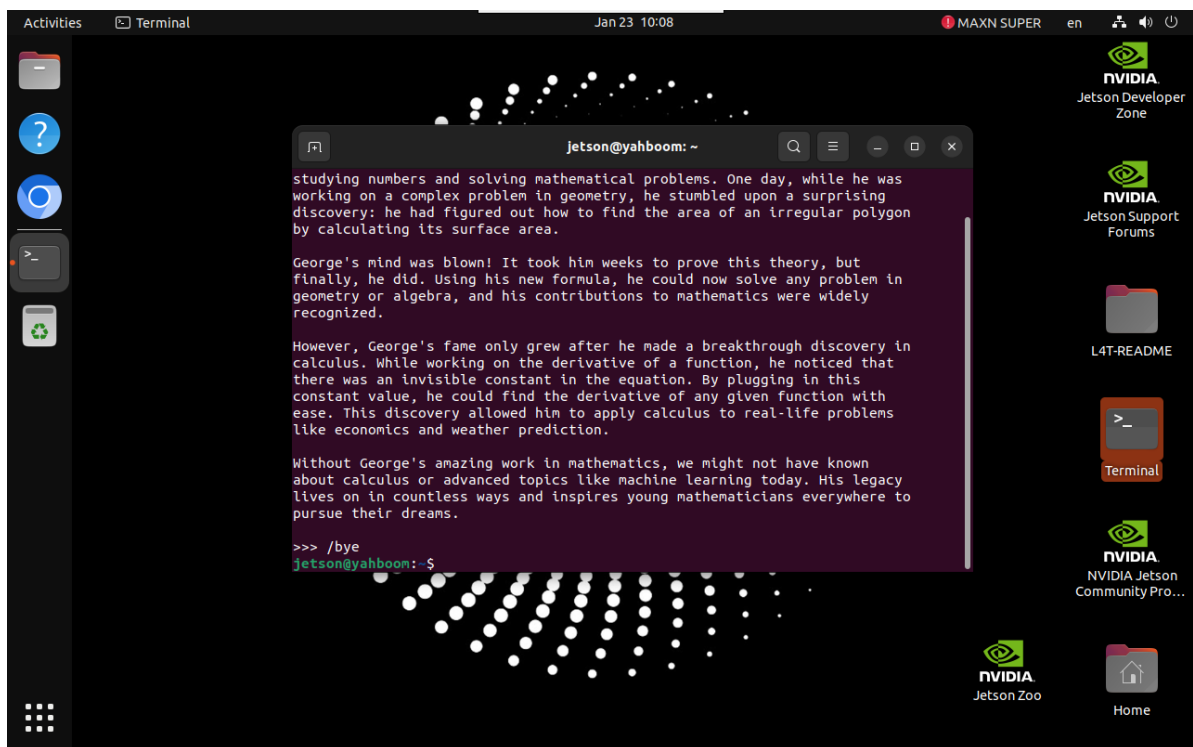
```
Tell me a mathematician story
```

The time to reply to the question is related to the hardware configuration, please be patient!



3.3. End the conversation

Use the `Ctrl+d` shortcut key or `/bye` to end the conversation!



References

Ollama

Official website: <https://ollama.com/>

GitHub: <https://github.com/ollama/ollama>

TinyLlama

GitHub: <https://github.com/jzhang38/TinyLlama>

Ollama corresponding model: <https://ollama.com/library/tinyllama>