

Phi4-mini

Phi4-mini

- 1. Model Scale
- 2. Performance
- 3. Using Phi4-mini
 - 3.1 Running Phi4-mini
 - 3.2 Starting a Conversation
 - 3.3 Ending a Conversation
- References

Demonstration Environment

Development Board: Jetson Orin Series Board


Note: Due to performance limitations, this model cannot be run on the Jetson Orin Nano 4GB.

The Phi-4-mini-instruction is a lightweight, open model built on synthetic data and curated public websites, focusing on high-quality, inference-intensive data.

1. Model Scale

Model	Volume
phi4-mini:3.8b	2.5GB

2. Performance



3. Using Phi4-mini

3.1 Running Phi4-mini

Use the run command to start running the model. If the model has not been downloaded, it will automatically pull the model from the Ollama model library:

```
ollama run phi4-mini:3.8b
```

```
root@root:~$ ollama run phi4-mini:3.8b
pulling manifest
pulling 3c168af1dea0: 100% 2.5 GB
pulling 813f53fdc6e5: 100% 655 B
pulling fa8235e5b48f: 100% 1.1 KB
pulling 8c2539a423c4: 100% 411 B
verifying sha256 digest
writing manifest
success
>>> Send a message (/? for help)
```

3.2 Starting a Conversation

```
How many minutes is a quarter of an hour?
```

Response time depends on your hardware configuration. Please be patient!

```
>>> How many minutes is a quarter of an hour?
A standard hour consists of 60 minutes. Therefore, one-quarter (1/4) of an
hour would be:

(1/4) * 60 = 15

So, there are 15 minutes in a quarter-hour.

>>> Send a message (/? for help)
```

3.3 Ending a Conversation

Use the `Ctrl+D` shortcut or `/bye` to end a conversation!

References

Ollama

Official Website: <https://ollama.com/>

GitHub: <https://github.com/ollama/ollama>

Phi4-mini

Ollama Compatible Model: <https://ollama.com/library/phi4-mini>