OpenClaw System Prompts

1. Course Content

Course Overview

- 💡 OpenClaw adopts an innovative **"file-as-prompt"** architecture — the AI Agent's behavior, identity, memory, tool usage rules, etc., are all defined and loaded through multiple Markdown files in the Workspace.

---

2. 🗁 Switching Workspace

- OpenClaw's system prompts and skills are stored in the workspace path: `$HOME/.openclaw/workspace`
- Jetson Orin NX comes with factory-preset workspace content in both Chinese and English, with identical content differing only in language
- Users in different regions can switch the workspace document language as needed
- Chinese workspace (prompts and skills content in Chinese)

```
cd /home/jetson/Language
sudo cp -a workspace_zh/. ~/.openclaw/workspace/
```

- English workspace (prompts and skills content in English)

```
cd /home/jetson/Language
sudo cp -a workspace_en/. ~/.openclaw/workspace/
```

- After configuration is complete, restart the OpenClaw gateway to apply the configuration

```
openclaw gateway restart
```

> [!WARNING]
>
> - The `.workspace_en` and `.workspace_zh` hidden folders under the `$HOME/.openclaw/workspace` path are Chinese and English template files. It is recommended not to modify them.

3. AGENTS.md

- Function Description

AGENTS.md defines the top-level specification for how an Agent operates in the workspace. 📖 Think of it as the Agent's **"user manual"** — telling the Agent how to start, how to manage memory, what it can do, and what it absolutely cannot do.

- Content Structure

AGENTS.md mainly contains the following core sections:

Section	Description
First Run	Guides the Agent's initialization process on first startup (reads BOOTSTRAP.md)
Mechanical Body	Describes the Agent's physical body (robotic arm, chassis, camera) and control methods
Session Startup	Specifies the startup process and file reading order for each session
Memory System	Defines the layered architecture for daily records and long-term memory
Red Line Rules	Lists absolutely prohibited behaviors (destructive commands, modifying robot program code)
Operating Standards	Distinguishes between operations that can be freely executed and those that require asking first
Memory Maintenance	Memory organization and update mechanisms during heartbeat intervals

> [!NOTE]
>
> ⚠ The factory AGENTS.md prompt restricts OpenClaw from modifying robot program code.

4. SOUL.md

- Function Description

SOUL.md is the Agent's **"soul" and core** 🧠. It defines the Agent's thinking approach and behavioral guidelines when facing tasks — how to control the robot body, how to find and use skills, how to plan and execute actions, and under what circumstances it should pause or refuse to execute. SOUL.md shapes the Agent's core behavioral pattern as a "robot operator."

- Working Principle

SOUL.md is read by the Agent at session startup as the underlying guiding principle for task execution. It does not specify specific skill parameters, but rather defines **the thinking process the Agent should follow when encountering any task**. For example, the Agent is required to always check the skill library before starting any task and cannot execute based on feeling or memory alone — this "check documentation first, then act" habit is defined by SOUL.md.

5. IDENTITY.md

- Function Description

IDENTITY.md defines the AI Agent's **self-awareness and personality traits**. It defines who the Agent "is" — its name, personality, language style, and basic identity positioning. This file gives the Agent a consistent personality when interacting with users.

- Working Principle

At each session startup, the Agent reads IDENTITY.md and establishes its "self-awareness" based on this file. This information affects the Agent's tone when answering questions, the way it proactively offers help, and the style of building relationships with users. For example, an Agent set to be "friendly and gentle" will naturally use a more approachable language style when replying to users.

6. HEARTBEAT.md

- Function Description

HEARTBEAT.md is the configuration file that controls the **heartbeat detection mechanism** in OpenClaw ⏰. The heartbeat mechanism allows the Agent to periodically "wake up" and perform background checks and maintenance tasks even without the user actively starting a session. This is similar to a robot's automatic scheduled inspection — the Agent can silently complete memory organization, status checks, and other work in the background.

- Working Principle

The OpenClaw gateway periodically checks the contents of the HEARTBEAT.md file to determine whether it needs to trigger the Agent's heartbeat API call:

- **File is empty or contains only comments**: Skip the heartbeat call; the Agent will not be woken up in the background

> [!NOTE]
>
> 💤 In the robot's factory configuration, the heartbeat mechanism is disabled by default.

---

7. 🛠️ TOOLS.md

- Function Description

TOOLS.md is the Agent's **tool usage localization notes** 📝. In the OpenClaw system, "Skills" define the general working methods of tools, while TOOLS.md records the Agent's own specific configurations and personalized settings. In short: Skills are the manuals (universal), and TOOLS.md is your personal notes (unique).

- Working Principle

During a session, the Agent reads both the relevant Skills documentation to understand the general usage of tools and TOOLS.md to obtain specific configurations related to its deployment environment.

8. 👤 USER.md

- Function Description

USER.md is the dedicated file used by the Agent to **understand and remember user information**. If IDENTITY.md is the window through which the Agent understands "who I am," then USER.md is the window through which the Agent understands "who it is dealing with." It records the user's basic information, preferences, project background, and other contextual information, helping the Agent provide more personalized and thoughtful service.

- Working Principle

At each session startup, the Agent reads USER.md according to the process specified in AGENTS.md. Through this file, the Agent can:

- Know how to address the user (name, form of address)
- Understand the user's timezone, preferences, and other background information
- Remember the user's concerns and ongoing projects

9. 🧠 MEMORY.md

- Function Description

MEMORY.md is the AI Agent's **long-term memory file** 🗂. In OpenClaw's memory system, there are two levels of memory storage:

1. 📓 **Daily Records** (`memory/YYYY-MM-DD.md`): Records raw events and conversation logs that occur each day, similar to a human "diary"
2. 🗃 **Long-term Memory** (`MEMORY.md`): Essential information distilled from daily records, similar to human "long-term memory" — containing important events, lessons learned, key decisions, and insights

- Working Principle

The loading and usage of MEMORY.md follow these rules:

- **Loaded only in the main session**: When the user communicates directly with the Agent, the Agent reads MEMORY.md to obtain long-term memory
- **Maintained autonomously by the Agent**: The Agent can freely read, edit, and update MEMORY.md in the main session

> [!NOTE]
>
> In actual use, users can:

- ✏️ **Actively write important information**: Directly write content you think the Agent should remember long-term into MEMORY.md
- 👀 **Observe the Agent's autonomous maintenance**: After the Agent has been running for a while, view the content automatically accumulated in MEMORY.md
- 🔄 **Manage memory content**: Regularly check and organize MEMORY.md, remove outdated information, and ensure long-term memory remains concise and accurate